

# Stock Forecasting Project

## Libraries

```
library(tidyverse)
library(lubridate)
library(quantmod)
library(tseries)
library(forecast)
```

## Data Gathering

### User Inputs (Stock and Date)

```
symbol <- "AAPL"
start_date <- as.Date("2018-01-01")
end_date <- Sys.Date()
```

### Stock Data Collection

```
getSymbols(symbol,
  src = "yahoo",
  from = start_date,
  to = end_date,
  auto.assign = TRUE)
```

```
## [1] "AAPL"
```

```
stock_data <- get(symbol)
```

```
head(stock_data)
```

```
##           AAPL.Open AAPL.High AAPL.Low AAPL.Close AAPL.Volume AAPL.Adjusted
## 2018-01-02   42.5400   43.0750  42.3150   43.0650   102223600   40.47983
## 2018-01-03   43.1325   43.6375  42.9900   43.0575   118071600   40.47280
## 2018-01-04   43.1350   43.3675  43.0200   43.2575    89738400   40.66079
## 2018-01-05   43.3600   43.8425  43.2625   43.7500   94640000   41.12372
## 2018-01-08   43.5875   43.9025  43.4825   43.5875   82271200   40.97097
## 2018-01-09   43.6375   43.7650  43.3525   43.5825   86336000   40.96627
```

## Data Exploration & Feature Engineering

### EDA

```
# Convert time-series data (xts object) to a regular tibble
df_stock <- tibble(
```

```

date      = index(stock_data),
open      = as.numeric(stock_data[, paste0(symbol, ".Open")]),
high      = as.numeric(stock_data[, paste0(symbol, ".High")]),
low       = as.numeric(stock_data[, paste0(symbol, ".Low")]),
close     = as.numeric(stock_data[, paste0(symbol, ".Close")]),
volume    = as.numeric(stock_data[, paste0(symbol, ".Volume")]),
adjusted  = as.numeric(stock_data[, paste0(symbol, ".Adjusted")])
)

glimpse(df_stock)

## Rows: 1,811
## Columns: 7
## $ date      <date> 2018-01-02, 2018-01-03, 2018-01-04, 2018-01-05, 2018-01-08, ~
## $ open      <dbl> 42.5400, 43.1325, 43.1350, 43.3600, 43.5875, 43.6375, 43.2900~
## $ high      <dbl> 43.0750, 43.6375, 43.3675, 43.8425, 43.9025, 43.7650, 43.5750~
## $ low       <dbl> 42.3150, 42.9900, 43.0200, 43.2625, 43.4825, 43.3525, 43.2500~
## $ close     <dbl> 43.0650, 43.0575, 43.2575, 43.7500, 43.5875, 43.5825, 43.5725~
## $ volume    <dbl> 102223600, 118071600, 89738400, 94640000, 82271200, 86336000,~
## $ adjusted  <dbl> 40.47983, 40.47280, 40.66079, 41.12372, 40.97097, 40.96627, 4~

summary(df_stock)

##           date           open           high           low
## Min.      :2018-01-02   Min.      : 35.99   Min.      : 36.43   Min.      : 35.50
## 1st Qu.: 2019-10-19   1st Qu.: 58.93   1st Qu.: 59.47   1st Qu.: 58.57
## Median : 2021-08-06   Median :136.82   Median :138.59   Median :134.92
## Mean     : 2021-08-07   Mean     :127.20   Mean     :128.59   Mean     :125.89
## 3rd Qu.: 2023-05-24   3rd Qu.:172.74   3rd Qu.:174.18   3rd Qu.:171.19
## Max.     : 2025-03-17   Max.     :258.19   Max.     :260.10   Max.     :257.63
##           close        volume        adjusted
## Min.      : 35.55   Min.      : 23234700   Min.      : 33.92
## 1st Qu.: 59.01   1st Qu.: 60953250   1st Qu.: 56.98
## Median :136.96   Median : 85671900   Median :134.38
## Mean     :127.31   Mean     : 98860006   Mean     :125.35
## 3rd Qu.:173.00   3rd Qu.:119555100   3rd Qu.:171.49
## Max.     :259.02   Max.     :426510000   Max.     :258.74

# Plot Adjusted Closing Price over time with a dynamic title
ggplot(df_stock, aes(x = date, y = adjusted)) +
  geom_line() +
  labs(title = paste(symbol, "Adjusted Closing Price"),
       x = "Date",
       y = "Adjusted Price") +
  theme_minimal()

```

## AAPL Adjusted Closing Price



```
# Check if there is any NA (rare to have NA)
df_stock %>%
  summarize(across(everything(), ~ sum(is.na(.))))
```

```
## # A tibble: 1 x 7
##   date    open  high    low close volume adjusted
##   <int> <int> <int> <int> <int>   <int>   <int>
## 1      0      0      0      0      0       0       0
```

## Feature Engineering

Add new variables

```
# Calculate Bollinger Bands using high, low, and close
bb <- BBands(HLC = df_stock %>% select(high, low, close),
             n = 20, maType = "SMA", sd = 2)

# Calculate MACD based on the adjusted closing price
macd_values <- MACD(df_stock$adjusted, nFast = 12, nSlow = 26, nSig = 9)

# Add all new variables
df_stock <- df_stock %>%
  arrange(date) %>%
  mutate(
    daily_return = (adjusted - lag(adjusted)) / lag(adjusted) * 100,
    lag1_close = lag(adjusted, 1),
```

```

lag2_close = lag(adjusted, 2),
ma20 = rollmean(adjusted, k = 20, fill = NA, align = "right"),
ma50 = rollmean(adjusted, k = 50, fill = NA, align = "right"),
rsi14 = RSI(adjusted, n = 14),
bb_dn   = bb[, "dn"],
bb_mavg = bb[, "mavg"],
bb_up   = bb[, "up"],
bb_pctB = bb[, "pctB"],
macd    = macd_values[, "macd"],
macdSig = macd_values[, "signal"],
rolling_sd_20 = rollapply(daily_return, width = 20,
                          FUN = sd, fill = NA, align = "right"),

wday = wday(date, label = TRUE),
sin_wday = sin(2 * pi * wday(date) / 7),
cos_wday = cos(2 * pi * wday(date) / 7)
)


```

### Visualize new variables

```


# Plot 14-Day RSI
ggplot(df_stock, aes(x = date, y = rsi14)) +
  geom_line() +
  labs(title = "14-Day RSI", x = "Date", y = "RSI") +
  theme_minimal()

```




```
# Plot Adjusted Price vs. Moving Averages using dynamic title
df_stock_long <- df_stock %>%
  select(date, adjusted, ma20, ma50) %>%
  pivot_longer(cols = c("adjusted", "ma20", "ma50"),
    names_to = "variable", values_to = "value")

ggplot(df_stock_long, aes(x = date, y = value, color = variable)) +
  geom_line() +
  labs(title = paste(symbol, "Adjusted Price vs MAs"),
    x = "Date", y = "Value", color = "Series") +
  theme_minimal()
```




```
# Plot Bollinger Bands with the closing price
ggplot(df_stock, aes(x = date)) +
  geom_line(aes(y = close), color = "blue") +
  geom_line(aes(y = bb_dn), color = "red", linetype = "dashed") +
  geom_line(aes(y = bb_up), color = "red", linetype = "dashed") +
  labs(title = "Bollinger Bands", y = "Price") +
  theme_minimal()
```

## Bollinger Bands



```
# Plot 20-day Rolling Std Dev of Daily Returns
ggplot(df_stock, aes(x = date, y = rolling_sd_20)) +
  geom_line() +
  labs(title = "20-day Rolling Std Dev of Daily Returns",
       x = "Date", y = "Rolling SD (%)") +
  theme_minimal()
```

## 20-day Rolling Std Dev of Daily Returns



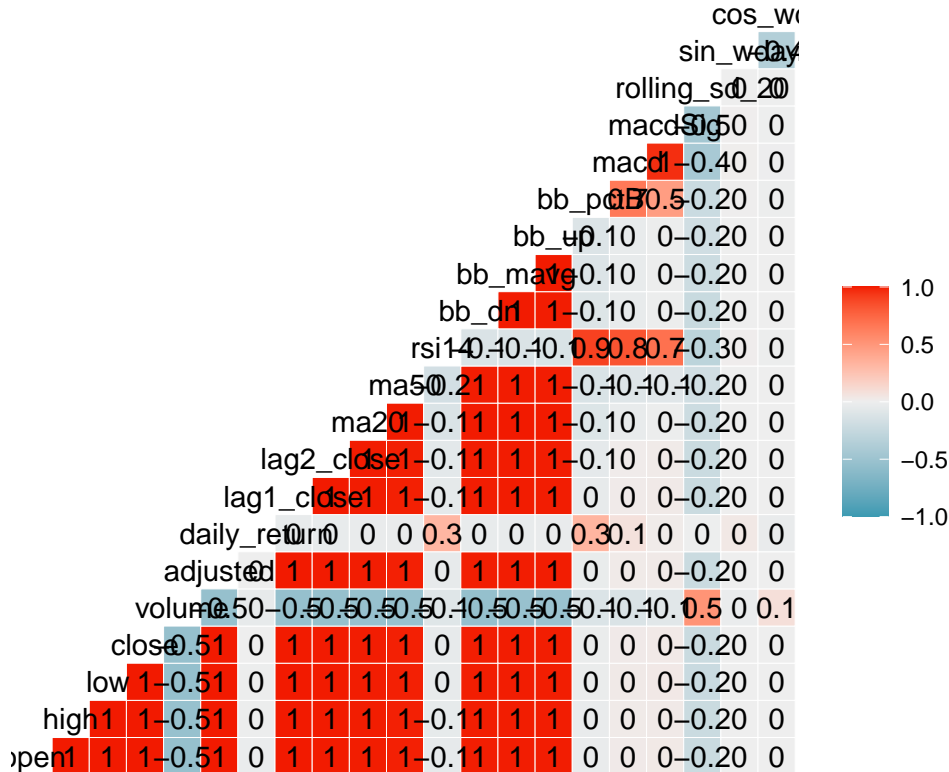
### Other tests and analysis

```
# Correlation Matrix of Numeric Features
df_numerics <- df_stock %>%
  select(where(is.numeric)) %>%
  drop_na() # Remove rows with NA for accurate correlation

GGally::ggcorr(df_numerics,
               method = c("pairwise.complete.obs", "pearson"),
               label = TRUE) +
  ggtitle("Correlation Matrix of Numeric Features")

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

## Correlation Matrix of Numeric Features



*# Stationarity Tests: Adjusted Price and Daily Returns*

```
adf_result <- adf.test(df_stock$adjusted, alternative = "stationary")
print(adf_result)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: df_stock$adjusted
## Dickey-Fuller = -3.2358, Lag order = 12, p-value = 0.08174
## alternative hypothesis: stationary
```

```
adf_returns <- adf.test(na.omit(df_stock$daily_return), alternative = "stationary")
```

```
## Warning in adf.test(na.omit(df_stock$daily_return), alternative =
## "stationary"): p-value smaller than printed p-value
```

```
print(adf_returns)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: na.omit(df_stock$daily_return)
## Dickey-Fuller = -11.491, Lag order = 12, p-value = 0.01
## alternative hypothesis: stationary
```

## Model 1: ARIMA with Exogenous Regressors (ARIMAX)